

html_notebook: default

Packages & Libraries

```
```{r} library(corrplot) library(ggplot2) library(caret) library(DMwR) library(mlbench) library(ROCR)
```

```
Preprocessing
```

```
```{r}
data <- read.csv("bank-additional-full.csv", sep=";")
head(data)
```

showing row count ```{r} cat("Original row count is", nrow(data))

```
dropping 'duration', 'pdays', 'default'
```{r}
data <- subset(data, select = -c(duration, pdays, default))
colnames(data)
```

find the columns that have "unknown" as a value ```{r} has\_unknown <- function(df, col) { (sum(df[,col]=="unknown")>0) }  
where\_is\_unknown <- function(df) { c <- 0 for(col in colnames(df)) { if (has\_unknown(df,col)) { cat(col) c <- c + 1 } } if(c == 0) {  
print(paste("There are no unknown values in the dataframe")) } } cat("These columns have unknown values")  
where\_is\_unknown(data)

```
replace those "unknown" values with a sample of the other values in those columns
```{r}
replace_unknowns <- function(df) {
  for(col in colnames(df)) {
    if(has_unknown(df,col)) {
      n_unk <- sum(df[,col]=="unknown")
      idx <- which(df[,col]=="unknown")
      df[idx,col] <- sample(col[!col=="unknown"], n_unk, replace=TRUE)
    }
  }
  df
}
data <- replace_unknowns(data)
head(data)
```

just to make sure those values were actually replaced and now there's no more unknowns ```{r} cat("# of unknowns:",
sum(data=="unknown"))

```
show the categorical columns
```{r}
cats <- names(data)[sapply(data,is.character)]
cats
```

encode all of the categorical values into numerals ```{r} encode <- function(df,col) { as.numeric(factor(df[,col]))-1 } for(cat in cats) { data[,cat] <- encode(data,cat) } head(data)

```
count plot for target to check for imbalance
```{r}
plt <- ggplot(data, aes(x = y)) + geom_bar(size=6)
plt + ggtitle("\t\t\t\t\tClass Imbalance in y")
```

Fixing the class imbalance by oversampling & resetting dataframe ```{r} data\$y <- factor(data\$y) data <- upSample(x=data[,ncol(data)],y=data\$y) names(data)[names(data)=='Class'] <- 'y'

```
count plot for target to ensure classes have been balanced
```{r}
plt <- ggplot(data, aes(x = y)) + geom_bar(size=6)
plt + ggtitle("\t\t\t\t\tClass Balance in y")
```

checking for NA's ```{r} print(paste("# of NA's", sum(is.na(data))))

```
checking for unknowns
```{r}
where_is_unknown(data)
```

since there are no unknown values and no NA's, lastly checking for increase in row count in dataframe ```{r} cat("New row count is",nrow(data))

```
#### Train Test Split

splitting 80/20 into train and test
```{r}
set.seed(3)
i <- sample(1:nrow(data), nrow(data)*0.80, replace = FALSE)
train <- data[i,]
test <- data[-i,]
dim(train)
dim(test)
```

## General EDA

plotting correlation matrix ```{r} train\$y <- encode(train,'y') C <- cor(train) corrplot(C, method='color', order = 'alphabet')

```

looking more specifically relative to 'y'
```{r}
C[, 'y']

```

it looks like euribor3m, nr.employed, emp.var.rate have the strongest negative correlation with y, and previous and poutcome, have the strongest positive correlation.

Linear Regression

to explore more, making base model for all independents against 'y' ```{r} set.seed(10) base <- lm(y~.,data=train) summary(base)

```

the insignificants appear to be age, job, housing, and loan, removing these and redoing model
```{r}
set.seed(11)
m1 <- lm(y~.,data=subset(train,select = -c(age,job,housing,loan)))
summary(m1)

```

every predictor is significant to the regression, but still  $R^2$  is low. let's see if scaling the independents help at all ```{r} minMaxScaler <- function(v) { (v-min(v))/(max(v)-min(v)) }

```

```{r}
train_scaled <- subset(train,select = -c(age,job,housing,loan))
train_scaled[, -ncol(train_scaled)] <- lapply(train_scaled[, -ncol(train_scaled)],minMaxScaler)
head(train_scaled)

```

```{r} set.seed(13) m2 <- lm(y~.,data=train\_scaled) summary(m2)

```

```{r}
par(mfrow=c(2,2))
plot(m2)
par(mfrow=c(1,1))

```

true to form, scaling does not impact the regression results. Both R^2 and the adjusted R^2 are staying low. we will use the model 'm1' for prediction as 'm2' doesn't appear to be doing much better.

Linear Regression Evaluation

create model for 'm1' once again (no scaling) ```{r} set.seed(11) m1 <- lm(y~.,data=subset(train,select = -c(age,job,housing,loan))) summary(m1)

```

get predictions and measure accuracy
```{r}
set.seed(21)
probas <- predict(m1, newdata = test, type="response")
pred <- ifelse(probas>0.5, 1, 0)
acc <- mean(pred==test$y)
cat("Accuracy = ", round(acc,3))
cat("Correlation = ",round(cor(probas,as.numeric(test$y)-1),3))

```

construct confusion matrix for easy visualization ```{r} confusionMatrix(as.factor(pred), reference=test\$y)

```
construct ROC plot & show auc
```{r}
set.seed(31)
pred <- prediction(probas,test$y)
roc <- performance(pred, measure="tpr", x.measure="fpr")
plot(roc)
cat("auc =", round(performance(pred,measure="auc")@y.values[[1]],3))
```

Overall, the linear regression model was predicting at around 73% accuracy. The correlation is around 51%. The model seemed to be able to distinguish between positive and negative classes for the 'y' feature though, with auc of about 78%.

Let's see if KNN regression can do more justice.

KNN Regression

For KNN we will use the scaled version of the training sample. We will need to scale everything in test as well except for the the target 'y'

```
```{r} test_scaled <- subset(test,select = -c(age,job,housing,loan)) test_scaled[,-ncol(test_scaled)] <- lapply(test_scaled[, -
ncol(test_scaled)],minMaxScaler)
```

```
Fit the scaled training data to the model

```{r}
kr1 <- knnreg(train_scaled[,-ncol(train_scaled)], train_scaled[, 'y'])
pred <- predict(kr1, test_scaled[,-ncol(test_scaled)])
cor_kr1 <- cor(pred,as.numeric(test$y)-1)
cat("correlation=", round(cor_kr1,3))
```

Correlation went up, let's see if we can select the best k to maximize it & minimize mse

```
{r} cor_kr <- rep(0,40) mse_kr <- rep(0,40) test_y_num <- encode(test,'y') i <- 1 for(k in seq(1,39,1))
{ fit_k <- knnreg(train_scaled[,-ncol(train_scaled)], train_scaled$y, k=k) pred_k <- predict(fit_k,
test_scaled[,-ncol(test_scaled)]) cor_kr[i] <- cor(pred_k,test_y_num) mse_kr[i] <- mean((pred_k -
test_y_num)^2) print(paste("k = ", k, cor_kr[i], mse_kr[i])) i <- i+1 }
```

While k=1 produces the largest correlation, k = 3, produces the smallest mse. See that k = 3, produces the second highest correlation. We will use k = 3 to train again.